

Where the acceleration goes: compute-bound vs branchy workloads

